

```
In [1]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

#ignore warning
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv(r"D:\Data Science with AI\2nd,3rd-jan-2024\Project\adult.csv\
```

```
In [3]: df
```

Out[3]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family
2	66	?	186061	Some-college	10	Widowed	?	Unmarried
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child
...
32556	22	Private	310152	Some-college	10	Never-married	Protective-serv	Not-in-family
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband
32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried
32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child

32561 rows × 15 columns



check dataset and shape

```
In [4]: df.shape
```

```
Out[4]: (32561, 15)
```

```
In [5]: df.head()
```

```
Out[5]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	r
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	W
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	W
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	B
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	W
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	W

```
In [6]: df.columns
```

```
Out[6]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education.num',  
              'marital.status', 'occupation', 'relationship', 'race', 'sex',  
              'capital.gain', 'capital.loss', 'hours.per.week', 'native.country',  
              'income'],  
            dtype='object')
```

view summary dataframe

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32561 entries, 0 to 32560  
Data columns (total 15 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   age                   32561 non-null  int64  
1   workclass             32561 non-null  object  
2   fnlwgt                32561 non-null  int64  
3   education             32561 non-null  object  
4   education.num         32561 non-null  int64  
5   marital.status        32561 non-null  object  
6   occupation            32561 non-null  object  
7   relationship          32561 non-null  object  
8   race                  32561 non-null  object  
9   sex                   32561 non-null  object  
10  capital.gain          32561 non-null  int64  
11  capital.loss          32561 non-null  int64  
12  hours.per.week        32561 non-null  int64  
13  native.country        32561 non-null  object  
14  income                32561 non-null  object
```

Encode ? as NANS

```
In [8]: df[df == '?'] = np.nan
```

again check summary dataframe

```
In [9]: df.shape
```

```
Out[9]: (32561, 15)
```

impute missing values with mode

```
In [10]: for col in ['workclass', 'occupation', 'native.country']:  
         df[col].fillna(df[col].mode()[0], inplace=True)
```

check again missing values

```
In [11]: df.isnull().sum()
```

```
Out[11]: age                0  
workclass                0  
fnlwgt                  0  
education                0  
education.num           0  
marital.status          0  
occupation              0  
relationship            0  
race                    0  
sex                     0  
capital.gain            0  
capital.loss            0  
hours.per.week          0  
native.country          0  
income                  0  
dtype: int64
```

setting feature vector and target variable

```
In [12]: x=df.drop(['income'],axis=1)  
         y=df['income']
```

```
In [13]: x.head
```

```
Out[13]: <bound method NDFrame.head of
cation.num      marital.status  \
0      90      Private    77053      HS-grad      9      Widowed
1      82      Private   132870      HS-grad      9      Widowed
2      66      Private   186061  Some-college     10      Widowed
3      54      Private   140359      7th-8th      4      Divorced
4      41      Private   264663  Some-college     10      Separated
...      ...      ...      ...      ...      ...      ...
32556    22      Private   310152  Some-college     10      Never-married
32557    27      Private   257302    Assoc-acdm     12  Married-civ-spouse
32558    40      Private   154374      HS-grad      9  Married-civ-spouse
32559    58      Private   151910      HS-grad      9      Widowed
32560    22      Private   201490      HS-grad      9      Never-married

      occupation  relationship  race  sex  capital.gain  \
0      Prof-specialty  Not-in-family  White  Female      0
1      Exec-managerial  Not-in-family  White  Female      0
2      Prof-specialty      Unmarried  Black  Female      0
3      Machine-op-inspct      Unmarried  White  Female      0
4      Prof-specialty      Own-child  White  Female      0
...      ...      ...      ...      ...      ...
32556  Protective-serv  Not-in-family  White  Male      0
32557      Tech-support      Wife  White  Female      0
32558  Machine-op-inspct      Husband  White  Male      0
32559      Adm-clerical      Unmarried  White  Female      0
32560      Adm-clerical      Own-child  White  Male      0

      capital.loss  hours.per.week  native.country
0      4356      40  United-States
1      4356      18  United-States
2      4356      40  United-States
3      3900      40  United-States
4      3900      40  United-States
...      ...      ...      ...
32556      0      40  United-States
32557      0      38  United-States
32558      0      40  United-States
32559      0      40  United-States
32560      0      20  United-States

[32561 rows x 14 columns]>
```

split the data into training and testing set

```
In [14]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.20,random_stat
```

feature engineering

#encode categorical categorical variable

```
In [15]: from sklearn import preprocessing
categorical=['workclass', 'education', 'marital.status', 'occupation', 'relation']
for feature in categorical :
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])
```

feature scaling

```
In [16]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=pd.DataFrame(sc.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(sc.transform(x_test),columns=x.columns)
```

```
In [17]: x_train.head(10)
```

Out[17]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relation
0	1.039817	-2.817895	0.960050	1.212397	-0.041721	-1.745708	-1.534809	-0.279
1	-1.407472	-0.085865	-0.461075	1.212397	-0.041721	0.915041	0.746328	0.963
2	-0.327786	-0.085865	-0.192395	0.177344	-0.429845	-1.745708	-0.014051	1.584
3	-1.047576	-0.085865	2.042380	0.177344	-0.429845	-0.415333	-1.027890	-0.90
4	-0.471744	-0.085865	-0.785647	0.177344	-0.429845	0.915041	0.239409	-0.279
5	0.895859	1.735487	-0.613820	0.436107	1.510776	-0.415333	-0.774430	-0.90
6	-0.111848	-0.085865	-0.027784	0.177344	-0.429845	0.915041	1.253248	1.584
7	0.248047	-0.085865	0.439758	-0.857710	0.734528	-0.415333	-1.027890	-0.90
8	1.327733	-0.085865	-1.143935	0.177344	-0.429845	-0.415333	1.253248	-0.90
9	-1.407472	-0.085865	1.526412	-2.669054	-1.594218	0.915041	-0.014051	0.963

logestic regression model with all features

```
In [18]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
lg=LogisticRegression()
lg.fit(x_train,y_train)
y_pred=lg.predict(x_test)
print('Logistic Regression accuracy score with all the features: {0:0.4f}'.format(accuracy_score(y_test,y_pred)))
```

Logistic Regression accuracy score with all the features: 0.8220

Explian variance ratio

```
In [19]: from sklearn.decomposition import PCA
pca = PCA()
x_train = pca.fit_transform(x_train)
pca.explained_variance_ratio_
```

```
Out[19]: array([0.14913354, 0.10195351, 0.08211085, 0.0799947 , 0.07478586,
0.07344007, 0.06851712, 0.0666489 , 0.06371768, 0.06181611,
0.05989432, 0.04841688, 0.04200177, 0.02756869])
```

logestic regression with first 13 feature

```
In [20]: x=df.drop(['income','native.country'], axis=1)
y=df['income']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=
categorical = ['workclass', 'education', 'marital.status', 'occupation', 'rela
for feature in categorical:
    le=preprocessing.LabelEncoder()
    x_train[feature]=le.fit_transform(x_train[feature])
    x_test[feature]=le.transform(x_test[feature])

x_train=pd.DataFrame(sc.fit_transform(x_train),columns=x.columns)
x_test=pd.DataFrame(sc.transform(x_test),columns=x.columns)
logreg=LogisticRegression()
logreg.fit(x_train,y_train)
y_pred=logreg.predict(x_test)
```

```
In [21]: 1 #Logestic regression 12 feature
2 x=df.drop(['income','native.country', 'hours.per.week'],axis=1)
3 y=df['income']
4
5
6 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_st
7 categorical=['workclass', 'education', 'marital.status', 'occupation', 're
8 for feature in categorical:
9     le=preprocessing.LabelEncoder()
10    x_train[feature]=le.fit_transform(x_train[feature])
11    x_test[feature]=le.transform(x_test[feature])
12 x_train=pd.DataFrame(sc.fit_transform(x_train),columns=x.columns)
13
14 x_test=pd.DataFrame(sc.transform(x_test),columns=x.columns)
15
16 logreg=LogisticRegression()
17 logreg.fit(x_train,y_train)
18 y_pred=logreg.predict(x_test)
```

```

In [24]: 1 # logistic regression with first 11 feature
2
3 x=df.drop(['income','native.country','hours.per.week','capital.loss'],
4 y=df['income']
5
6 x_train,x_test,y_train,y_test=train_test_split(x,y, test_size=0.3, random_
7
8 categorical=['workclass','education','marital.status','occupation','re
9 for feature in categorical:
10     le=preprocessing.LabelEncoder()
11     x_train[feature]=le.fit_transform(x_train[feature])
12     x_test[feature]=le.transform(x_test[feature])
13
14 x_train=pd.DataFrame(sc.fit_transform(x_train), columns=x.columns)
15 x_test=pd.DataFrame(sc.transform(x_test), columns=x.columns)
16
17 logreg=LogisticRegression()
18 logreg.fit(x_train,y_train)
19 y_pred=logreg.predict(x_test)

```

```

In [32]: #select number of dimension

x=df.drop(['income'],axis=1)
y=df['income']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=

categorical=['workclass','education','marital.status','occupation','relati
for feature in categorical:

    le = preprocessing.LabelEncoder()
    x_train[feature] = le.fit_transform(x_train[feature])
    x_test[feature] = le.transform(x_test[feature])


x_train = pd.DataFrame(sc.fit_transform(x_train), columns = x.columns)

pca= PCA()
pca.fit(x_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
dim = np.argmax(cumsum >= 0.90) + 1
print('The number of dimensions required to preserve 90% of variance is',dim)

```

The number of dimensions required to preserve 90% of variance is 12

```
In [34]: 1 #plot explained variance ratio with number of dimensions
2 plt.figure(figsize=(8,6))
3 plt.plot(np.cumsum(pca.explained_variance_ratio_))
4 plt.xlim(0,14,1)
5 plt.xlabel('Number of components')
6 plt.ylabel('Cumulative explained variance')
7 plt.show()
```



```
In [ ]:
```